

# **Text Classification**

Dr. Nguyen Van Vinh

CS Department – UET, Hanoi VNU

# Is this spam?

**Subject: Important notice!**

**From:** Stanford University <newsforum@stanford.edu>

**Date:** October 28, 2011 12:34:16 PM PDT

**To:** undisclosed-recipients:;

---

Greats News!

You can now access the latest news by using the link below to login to Stanford University News Forum.

<http://www.123contactform.com/contact-form-StanfordNew1-236335.html>

Click on the above link to login for more information about this new exciting forum. You can also copy the above link to your browser bar and login for more information about the new services.

© Stanford University. All Rights Reserved.

# Positive or negative movie review?

- unbelievably disappointing
- Full of zany characters and richly applied satire, and some great plot twists
- this is the greatest screwball comedy ever filmed
- It was pathetic. The worst part about it was the boxing scenes.

# Why text classification?

**BÁOMÓI**  
trang thông tin điện tử

Nhập nội dung tìm kiếm

Nguyen Van Vinh ▾

NÓNG MỚI <sup>11</sup> VIDEO CHỦ ĐỀ # Năng lượng tích cực # Khám phá Việt Nam # Đại hội Đảng các cấp # Phòng chống COVID-19

XÃ HỘI THẾ GIỚI VĂN HÓA KINH TẾ GIÁO DỤC

Thời sự  
Giao thông  
Môi trường - Khí hậu

Nghệ thuật  
Ẩm thực  
Du lịch

Lao động - Việc làm  
Tài chính  
Chứng khoán  
Kinh doanh

Học bổng - Du học  
Đào tạo - Thi cử

THỂ THAO GIẢI TRÍ PHÁP LUẬT CÔNG NGHỆ KHOA HỌC

Bóng đá quốc tế  
Bóng đá Việt Nam  
Quần vợt

Âm nhạc  
Thời trang  
Điện ảnh - Truyền hình

An ninh - Trật tự  
Hình sự - Dân sự

CNTT - Viễn thông  
Thiết bị - Phần cứng

ĐỜI SỐNG XE CỘ NHÀ ĐẤT

DỊCH VỤ

# Content

- Text classification problem
- Feature extraction from text
- Naïve Bayes method
- Logistic Regression method
- Sentiment Analysis Casestudy

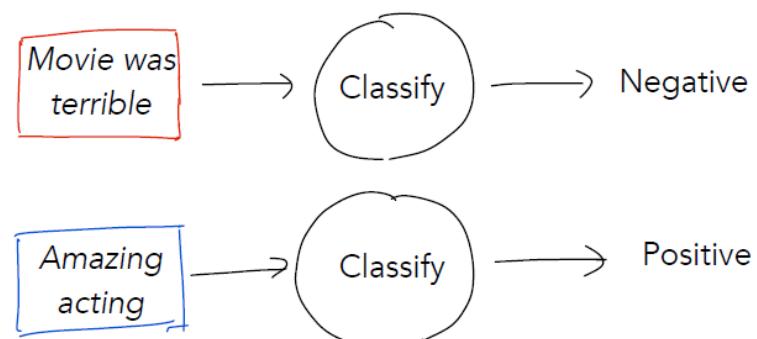
# Text Classification: definition

## Inputs:

- A document  $d$
- A set of classes  $C$  ( $m$  classes)

## Output:

- Predicted class  $c \in C$  for document  $d$



# Classification Methods:

## Hand-coded rules

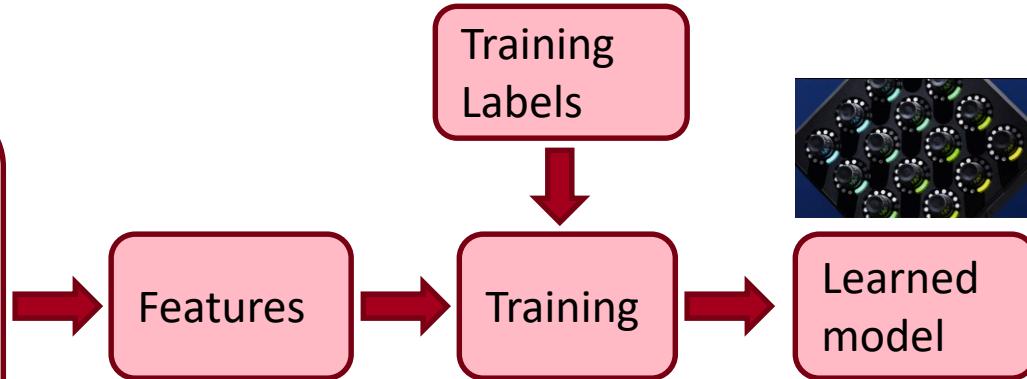
- Choose features
- Rules based on combinations of words or other features
  - spam: black-list-address OR (“dollars” AND “have been selected”)
- Accuracy can be high
  - If rules carefully refined by expert
- But building and maintaining these rules is expensive

# Classification Methods: Supervised Machine Learning

- *Input:*
    - a document  $d$
    - a fixed set of classes  $C = \{c_1, c_2, \dots, c_J\}$
    - A training set of  $m$  hand-labeled documents  $(d_1, c_1), \dots, (d_m, c_m)$   
 $d_i \subset D$
  - *Output:*
    - a learned classifier  $F: D \rightarrow C$
- Key questions:**
- a) What is the form of  $F$ ?
  - b) How do we learn  $F$ ?

# Machine Learning Architecture

## Learning



## Inference



# Classification Methods: Supervised Machine Learning

- Any kind of classifier
  - Naïve Bayes
  - Logistic regression
  - Support-vector machines
  - Random Forest/XGBoost/CatBoost/LightGBM
  - Neural Network
  - ...

# Precision and recall

- **Precision:** % of selected items that are correct  
**Recall:** % of correct items that are selected

	correct	not correct
selected	tp	fp
not selected	fn	tn

# A combined measure: F

- A combined measure that assesses the P/R tradeoff is F measure (weighted harmonic mean):

$$F = \frac{1}{\alpha \frac{1}{P} + (1-\alpha) \frac{1}{R}} = \frac{(\beta^2 + 1)PR}{\beta^2 P + R}$$

- The harmonic mean is a very conservative average; see *IIR* § 8.3
- People usually use balanced F1 measure
  - i.e., with  $\beta = 1$  (that is,  $\alpha = \frac{1}{2}$ ):

$$F = 2PR/(P+R)$$

# Features

- A measurable variable that is (rather, should be) distinctive of something we want to model.
- We usually choose features that are useful to **identify** something, i.e., to do **classification**
  - Ex: Cô gái đó **rất đẹp** trong bữa tiệc hôm đó.
- We often need **several** features to adequately model something – *but not too many!*

# Feature vectors

- Values for several features of an **observation** can be put into a single **vector**

 **Damien Fahey**   
@DamienFahey

Rush Limbaugh looks like if someone put a normal human being in landscape mode.

 Reply  Retweet  Favorite  More

 **Faux John Madden**  
@FauxJohnMadden

BREAKING: Apple Maps projecting Barack Obama to win Brazil.

 Reply  Retweet  Favorite  More

 **Jim Gaffigan**   
@JimGaffigan

If there was an award for most pessimistic, I probably wouldn't even be nominated.

 Reply  Retweet  Favorite  More

# proper nouns	# 1st person pronouns	# commas
2	0	0
5	0	0
0	1	1

# Feature vectors

- Features should be useful in **discriminating** between categories.

Table 3: Features to be computed for each text

- Counts:
  - First person pronouns
  - Second person pronouns
  - Third person pronouns
  - Coordinating conjunctions
  - Past-tense verbs
  - Future-tense verbs
  - Commas
  - Colons and semi-colons
  - Dashes
  - Parentheses
  - Ellipses
  - Common nouns
  - Proper nouns
  - Adverbs
  - wh*-words
  - Modern slang acronyms
  - Words all in upper case (at least 2 letters long)
- Average length of sentences (in tokens)
- Average length of tokens, excluding punctuation tokens (in characters)
- Number of sentences

Higher values → this person is referring to themselves (to their opinion, too?)

Higher values → looking forward to (or dreading) some future event?

Lower values → this tweet is more formal. Perhaps not overly sentimental?

# Classification: Sentiment Analysis

*this movie was great! would watch again* Positive

*that film was awful, I'll never watch again* Negative

- Surface cues can basically tell you what's going on here: presence or absence of certain words (*great, awful*)
- Steps to classification:
  - Turn examples like this into feature vectors
  - Pick a model / learning algorithm
  - Train weights on data to get our classifier

# Feature Representation

*this movie was great! would watch again* Positive

- Convert this example to a vector using *bag-of-words features*

[contains *the*] [contains *a*] [contains *was*] [contains *movie*] [contains *film*] ...  
position 0 position 1 position 2 position 3 position 4

$$f(x) = [0 \quad 0 \quad 1 \quad 1 \quad 0 \quad \dots]$$

- Very large vector space (size of vocabulary), sparse features
- Requires *indexing* the features (mapping them to axes)
- More sophisticated feature mappings possible (m-idf), as well as lots of other features: character n-grams, parts of speech, ...

# Feature Engineering for Text

- **Feature Engineering** is the process of creating new features or transforming existing features to improve the performance of a machine-learning model
- Tokenizing
- Bag of Words
- N-grams
- TF-IDF
- Embeddings
- NLP

# Feature Selection

- Which features to use?
- How many features to use?

Approaches:

- Frequency
- Mutual Information
- Accuracy

# Naive Bayes: Exercise!

- Model

$$P(x, y) = P(y) \prod_{i=1}^n P(x_i|y)$$

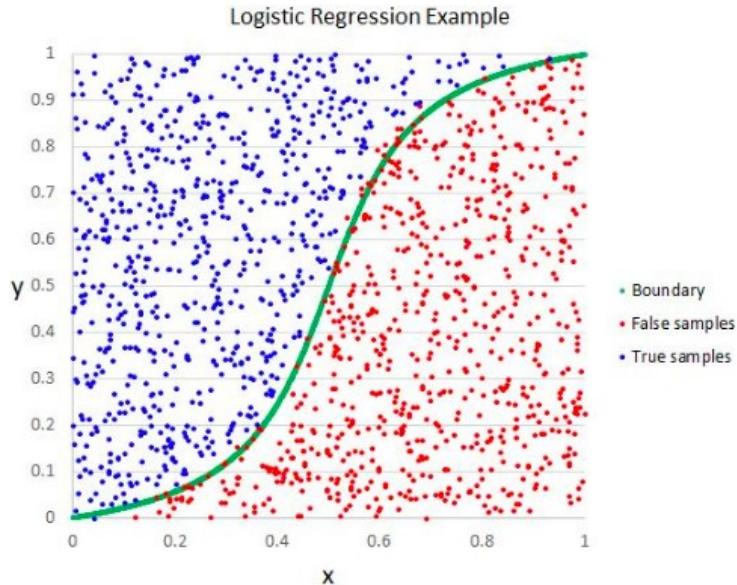
- Inference

$$\operatorname{argmax}_y \log P(y|x) = \operatorname{argmax}_y \left[ \log P(y) + \sum_{i=1}^n \log P(x_i|y) \right]$$

- Learning: maximize  $P(x,y)$  by reading counts off the data

# **LOGISTIC REGRESSION MODEL**

# Logistic Regression



- Powerful supervised model
- Baseline approach for many NLP tasks
- Is also the foundation of neural networks
- Binary (two classes) or multinomial (>2 classes)

# Generative and Discriminative Classifiers

- Naive Bayes is a **generative** classifier
- by contrast:
- Logistic regression is a **discriminative** classifier

# Generative and Discriminative Classifiers

Suppose we're distinguishing cat from dog images



imagenet



imagenet

# Generative Classifier:

- Build a model of what's in a cat image
  - Knows about whiskers, ears, eyes
  - Assigns a probability to any image:
    - how cat-y is this image?



Also build a model for dog images

Now given a new image:

**Run both models and see which one fits better**

# Discriminative Classifier

Just try to distinguish dogs from cats



Oh look, dogs have collars!  
Let's ignore everything else

# Finding the correct class $c$ from a document $d$ in Generative vs Discriminative Classifiers

- Naive Bayes

$$\hat{c} = \operatorname{argmax}_{c \in C} \underbrace{P(d|c)}_{\text{likelihood}} \underbrace{P(c)}_{\text{prior}}$$

- Logistic Regression

$$\hat{c} = \operatorname{argmax}_{c \in C} P(c|d) \quad \text{posterior}$$

# Overall process: Discriminative classifiers

**Input:** a set of labeled documents  $\{(d_i, y_i)\}_{i=1}^n$

- Components:

1. Convert  $d_i$  into a **feature representation**  $x_i$   
 $y_i = 0 \text{ or } 1$  (binary)  
 $y_i = 1, \dots, m$  (multinomial)
  2. **Classification function** to compute  $\hat{y}$  using  $P(\hat{y} | x)$   Using either sigmoid or softmax!
  3. **Loss function** for learning
  4. Optimization **algorithm**
- **Train phase:** Learn the **parameters** of the model to minimize **loss function** on the training set
  - **Test phase:** Apply **parameters** to predict class given a new input  $x$  (feature representation of testing document  $d$ )

# Feature representation

## Bag of words

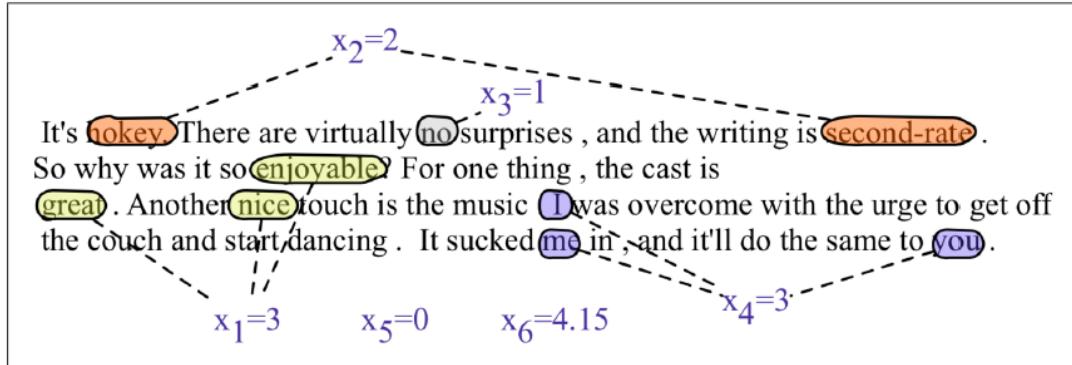
I love this movie! It's sweet, but with satirical humor. The dialogue is great and the adventure scenes are fun... It manages to be whimsical and romantic while laughing at the conventions of the fairy tale genre. I would recommend it to just about anyone. I've seen it several times, and I'm always happy to see it again whenever I have a friend who hasn't seen it yet!



$$\mathbf{x} = [x_1, x_2, \dots, x_k]$$

In BoW representations,  $k = |V|$  and the vector could be very sparse

# Example: Sentiment classification



Var	Definition	Value in Fig. 5.2
$x_1$	count(positive lexicon) $\in$ doc	3
$x_2$	count(negative lexicon) $\in$ doc	2
$x_3$	$\begin{cases} 1 & \text{if "no" } \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	1
$x_4$	count(1st and 2nd pronouns $\in$ doc)	3
$x_5$	$\begin{cases} 1 & \text{if "!" } \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	0
$x_6$	$\ln(\text{word count of doc})$	$\ln(64) = 4.15$

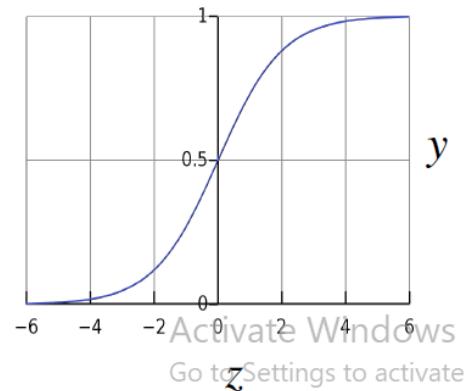
Remember that the values make up the feature vector!

Activate Windows  
Go to Settings to activat

# Classification function

- Given: Input feature vector  $\mathbf{x} = [x_1, x_2, \dots, x_k]$
- Output:  $P(y = 1 | \mathbf{x})$  and  $P(y = 0 | \mathbf{x})$  (binary classification)

- Weight vector  $\mathbf{w} = [w_1, w_2, \dots, w_k]$       bias
- Given input features  $\mathbf{x}$ :  $z = \mathbf{w} \cdot \mathbf{x} + b$
  - Therefore,  $\hat{y} = P(y = 1 | \mathbf{x}) = \sigma(\mathbf{w} \cdot \mathbf{x} + b) = \frac{1}{1 + e^{-(\mathbf{w} \cdot \mathbf{x} + b)}}$
  - Decision boundary:  $= \begin{cases} 1 & \text{if } \hat{y} > 0.5 \\ 0 & \text{otherwise} \end{cases}$



# Example: Sentiment classification

Var	Definition	Value
$x_1$	count(positive lexicon) $\in$ doc	3
$x_2$	count(negative lexicon) $\in$ doc	2
$x_3$	$\begin{cases} 1 & \text{if "no" } \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	1
$x_4$	count(1st and 2nd pronouns $\in$ doc)	3
$x_5$	$\begin{cases} 1 & \text{if "!" } \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	0
$x_6$	$\log(\text{word count of doc})$	$\ln(64) = 4.15$

- Assume weights  $\mathbf{w} = [2.5, -5.0, -1.2, 0.5, 2.0, 0.7]$  and bias  $b = 0.1$

$$\begin{aligned} p(+|x) &= P(Y = 1|x) = \sigma(w \cdot x + b) \\ &= \sigma([2.5, -5.0, -1.2, 0.5, 2.0, 0.7] \cdot [3, 2, 1, 3, 0, 4.15] + 0.1) \\ &= \sigma(.805) \\ &= 0.69 \end{aligned}$$

$$\begin{aligned} p(-|x) &= P(Y = 0|x) = 1 - \sigma(w \cdot x + b) \\ &= 0.31 \end{aligned}$$

Activate  
Go to Setti

# Idea of logistic regression

- We'll compute  $w \cdot x + b$
- And then we'll pass it through the sigmoid function:
  - $\sigma(w \cdot x + b)$
- And we'll just treat it as a probability

# Sigmoid Function

- Among the functions with the above 2 features, the function sigmoid:  $f(s) = \frac{1}{1+e^{-s}}$  widely used.
- $f'(s) = f(s) (1-f(s))$ .

# Loss function

- For n data points  $(x_i, y_i)$ ,  $\hat{y}_i = P(y_i = 1 \mid x_i)$
- Classifier probability:  $\prod_{i=1}^n P(y_i \mid x_i) = \prod_{i=1}^n \hat{y}_i^{y_i} (1 - \hat{y}_i)^{1-y_i}$
- Loss:  $-\log \prod_{i=1}^n P(y_i \mid x_i) = -\sum_{i=1}^n \log P(y_i \mid x_i)$

$$L_{CE} = - \sum_{i=1}^n [y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)]$$

# Properties of CE loss

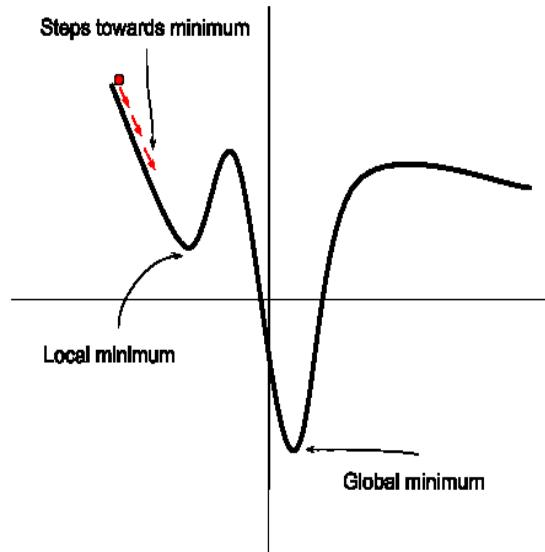
- $L_{CE} = - \sum_{i=1}^n [y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)]$
- What values can this loss take?
  - A) 0 to  $\infty$
  - B)  $-\infty$  to  $\infty$
  - C)  $-\infty$  to 0
  - D) 1 to  $\infty$

# Optimization of loss based on gradient decent

In more detail about Gradient decent:

<https://towardsdatascience.com/implement-gradient-descent-in-python-9b93ed7108d1>

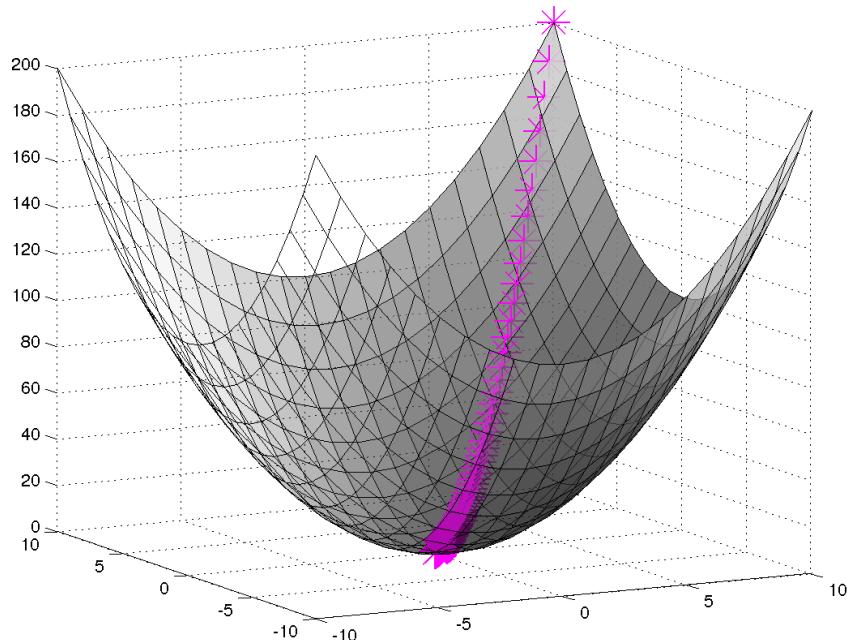
- It is an **optimization algorithm** to find the **minimum of a function**. We start with a random point on the function and move in the **negative direction of the gradient** of the function to reach the **local/global minima**.



# Gradient descent

- If the derivative of the function  $x$  at  $x_t$ :  $f'(x_t) > 0$  then  $x_t$  is to the right of  $x^*$  (and vice versa). In order for the next point  $x_{t+1}$  to be closer to  $x^*$ , we need to move  $x_t$  to the left, i.e negative side. In other words, we need to move **the opposite sign with the derivative** :  $x_{t+1} = x_t + \Delta$
- $x_{t+1} = x_t - \mu f'(x_t)$

# Gradient descent



# Gradient Descent Algorithm

---

**Algorithm 1** Gradient Descent

---

- 1: Initialize  $\theta_0 = k \in \mathbb{R}^d$
- 2: Initialize  $t \leftarrow 0$
- 3: **repeat**
- 4:    $\forall j \in \{1, \dots, d\}, \theta_{j,t+1} \leftarrow \theta_{j,t} - \alpha \frac{\partial}{\partial \theta_j} L(\theta) \Big|_{\theta=\theta_t}$
- 5:    $t \leftarrow t + 1$
- 6: **until** converge

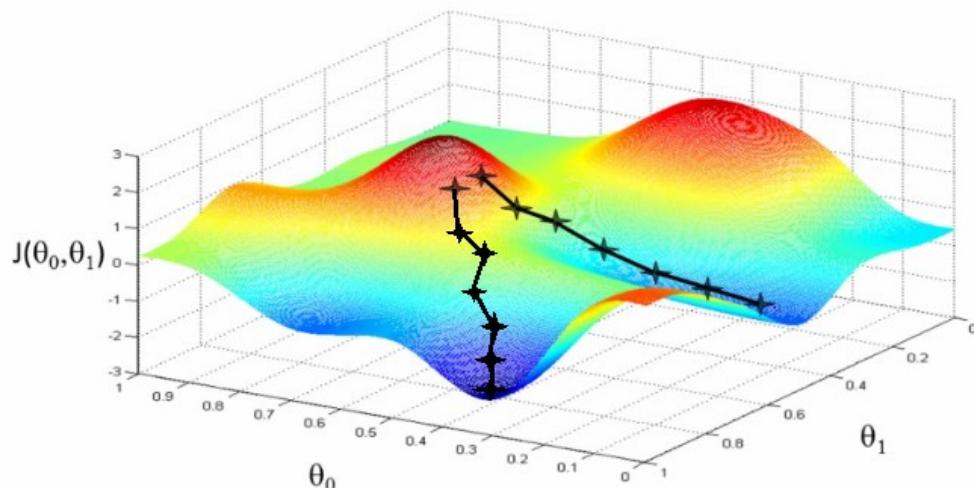
---

**Variants of GD algorithm:**

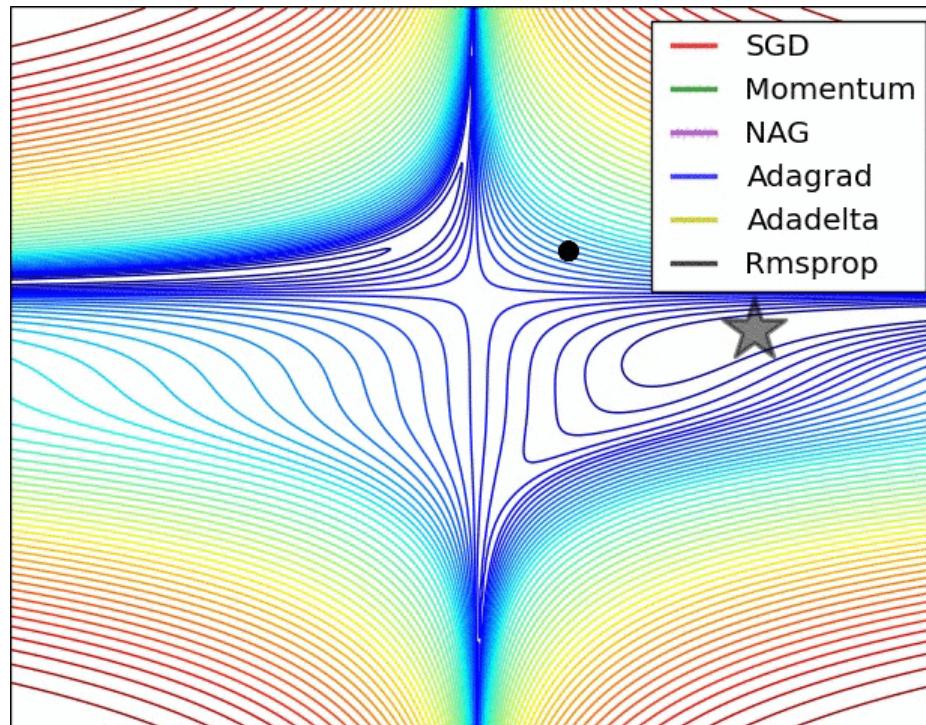
- Batch Gradient Descent
- Stochastic Gradient Descent
- Mini-batch Gradient Descent

# Pros and cons of gradient descent

- Simple and often quite effective on ML tasks
- Often very scalable
- Only applies to smooth functions (differentiable)
- Might find a local minimum, rather than a global one



# Visualization of algorithms of GD



# Gradient for logistic regression

$$\hat{y}_i = \sigma(\mathbf{w} \cdot \mathbf{x}_i + b)$$

$$L_{CE} = - \sum_{i=1}^n [y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)]$$

- Gradient,  $\frac{dL_{CE}(\mathbf{w}, b)}{dw_j} = \sum_{i=1}^n [\hat{y}_i - y_i] x_{i,j}$    
The j-th value of the feature vector  $\mathbf{x}_i$
- $\frac{dL_{CE}(\mathbf{w}, b)}{db} = \sum_{i=1}^n [\hat{y}_i - y_i]$

Cross entropy loss for logistic regression is **convex** (i.e. has only one global minimum) so gradient descent is guaranteed to find the minimum.

# Selecting learning rate

- Use grid-search in log-space over small values on a tuning set:
  - e.g., 0.01, 0.001, ...
- Sometimes, decrease after each pass:
  - e.g factor of  $1/(1 + dt)$ ,  $t=\text{epoch}$
  - sometimes  $1/t^2$
- Fancier techniques:
  - Adaptive gradient: scale gradient differently for each dimension  
(Adagrad, ADAM, ....)

# Logistic Regression: Summary

- ▶ Model

$$P(y = +|x) = \frac{\exp(\sum_{i=1}^n w_i x_i)}{1 + \exp(\sum_{i=1}^n w_i x_i)}$$

- ▶ Inference

$\operatorname{argmax}_y P(y|x)$  fundamentally same as Naive Bayes

$$P(y = 1|x) \geq 0.5 \Leftrightarrow w^\top x \geq 0$$

- ▶ Learning: gradient ascent on the (regularized) discriminative log-likelihood

# **SENTIMENT ANALYSIS**

# Google Product Search



**HP Officejet 6500A Plus e-All-in-One Color Ink-jet - Fax / copier / printer / scanner**

\$89 online, \$100 nearby    377 reviews

September 2010 - Printer - HP - Inkjet - Office - Copier - Color - Scanner - Fax - 250 sh

## Reviews

**Summary** - Based on 377 reviews



### What people are saying

ease of use		"This was very easy to setup to four computers."
value		"Appreciate good quality at a fair price."
setup		"Overall pretty easy setup."
customer service		"I DO like honest tech support people."
size		"Pretty Paper weight."
mode		"Photos were fair on the high quality mode."
colors		"Full color prints came out with great quality."

# Bing Shopping

## HP Officejet 6500A E710N Multifunction Printer

[Product summary](#) [Find best price](#) **Customer reviews** [Specifications](#) [Related items](#)



\$121.53 - \$242.39 (14 stores)

Compare

Average rating (144)

(55)

(54)

(10)

(6)

(23)

(0)

Most mentioned

Performance (57)

Ease of Use (43)

Print Speed (39)

Connectivity (31)

More ▾

Show reviews by source

Best Buy (140)

CNET (5)

Amazon.com (3)

# Sentiment Analysis task

- Input: Given a sentence (paragraph) of text.
- Output: Determine whether this sentence is positive (1) or negative (0).
- Movie review dataset: IMDB

# Sentiment Analysis

- Dataset: IMDB Movie Reviews
- Data preprocessing
- Features extraction (Vectorization)
- Choosing a machine learning model and training

# Data preprocessing

- The smallest unit is the word. String of words:

Ví dụ: **Text:** This is a cat. --> **Word Sequence:** [this, is, a, cat]

- Data crawled from the web is usually very “dirty” such as “Html”, abbreviations, ... should be careful with data preprocessing
- Using **Regular Expression (RE)**

# RE

- Regular Expression (or **regex**) is a sequence of characters that represent a search pattern
- Each character has a meaning; for example, “.” means any character that isn't the newline character: '\n'”
- These characters are often combined with quantifiers, such as \*, which means **zero or more**.
- **Regular expressions** are very useful when solving with string problem.

# RE

SYMBOL	USAGE
\$	Matches the end of the line
\s	Matches whitespace
\S	Matches any non-whitespace character
*	Repeats a character zero or more times
\S	Matches any non-whitespace character
*?	Repeats a character zero or more times (non-greedy)
+	Repeats a character one or more times
+?	Repeats a character one or more times (non-greedy)
[aeiou]	Matches a single character in the listed set
[^XYZ]	Matches a single character not in the listed set
[a-zA-Z0-9]	The set of characters can include a range

# Quiz

- Example:

- 1) Extract **number** from s:

s = 'My 2 favourite numbers are 8 and 25. My mobie is 0912203062'.

- 1) Extract **email** from s:

s = 'Hello from shubhamg199630@gmail.com to priya@yahoo.com about the meeting @2PM'

# Features extraction

- Now that we have a way to extract information from text in the form of word sequences, we need a way to transform these word sequences into numerical features: this is **vectorization**.
- The simplest vectorization technique is Bag Of Words (BOW). It starts with a list of words called the vocabulary (this is often all the words that occur in the training data)

# Features extraction

- To use BOW vectorization in Python, we can rely on CountVectorizer from the scikit-learn library
- scikit-learn has a built-in list of stop words that can be ignored by passing stop\_words="english" to the vectorizer
- Moreover, we can pass our custom pre-processing function from earlier to automatically clean the text before it's vectorized.

**Training texts:** ["This is a good cat", "This is a bad day"]=>

**vocabulary:** [this, cat, day, is, good, a, bad]

**New text:** "This day is a good day" --> [1, 0, 2, 1, 1, 1, 0]

# Sentiment Analysis CaseStudy

- Dữ liệu văn bản IMDB: The IMDB movie reviews dataset is a set of 50,000 reviews, half of which are positive and the other half negative
- We can download data from:  
[http://ai.stanford.edu/~amaas/data/sentiment/aclImdb\\_v1.tar.gz](http://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz)
- We have a directory containing data: aclImdb
- We can use the following function to load the training/test datasets from IMDB

# Sentiment Analysis Case Study

- feature vectors that result from BOW are usually very large (80,000-dimensional vectors in this case)
- we need to use simple algorithms that are efficient on a large number of features (e.g., Naive Bayes, linear SVM, or logistic regression)

# Improving to the current model

- Features extraction is very important (Features Engineering)
- There are some biases attached with only looking at how many times a word occurs in a text. In particular, the longer the text, the higher its features (word counts) will be
- Dựa vào đặc trưng TF-IDF
- Dựa vào n-gram

# Improving model

- TF-IDF feature
- We can train the Linear SVM model with the TF-IDF feature simply by replacing the function CountVectorizer by TfidfVectorizer
- Result improving about 2%

$$TF(\text{word}, \text{text}) = \frac{\text{number of times the word occurs in the text}}{\text{number of words in the text}}$$

$$IDF(\text{word}) = \log \left[ \frac{\text{number of texts}}{\text{number of texts where the word occurs}} \right]$$

$$TF-IDF(\text{word}, \text{text}) = TF(\text{word}, \text{text}) \times IDF(\text{word})$$

# Improving model

- Using independent words is not good. Ex:
- if the word **good** occurs in a text, we will naturally tend to say that this text is positive, even if the actual expression that occurs is actually **not good**. **Phrase is better.**
- Using n-gram in order to solve this problem
- An N-gram is a set of N successive words (e.g., very good [2-gram] and not good at all [4-gram]). Using N-grams, we produce richer word sequences.
- Ví dụ với N =2:

This is a cat. --> [this, is, a, cat, (this, is), (is, a), (a, cat)]

# Improving model

- In practice, including N-grams in our TF-IDF vectorizer is as simple as providing an additional parameter `ngram_range=(1, N)`.

```
vectorizer = TfidfVectorizer(stop_words="english",
                             preprocessor=clean_text,
                             ngram_range=(1, 2))
```

# Homework

- Sentiment Analysis with dataset IMDB (link:  
<http://ai.stanford.edu/~amaas/data/sentiment/>)

# References

- **Chapter 4,5** - Speech and Language Processing (3rd ed. draft)  
(<https://web.stanford.edu/~jurafsky/slp3/4.pdf>)  
(<https://web.stanford.edu/~jurafsky/slp3/5.pdf>)