

Word Embedding

Nguyen Van Vinh
VNU-UET

word2vec: The foundation of NLP

- <https://kharshit.github.io/blog/2018/07/27/word2vec-the-basic-of-nlp>


The important idea: model of meaning focusing on similarity

Each word = a vector

$$v_{\text{cat}} = \begin{pmatrix} -0.224 \\ 0.130 \\ -0.290 \\ 0.276 \end{pmatrix} \quad v_{\text{dog}} = \begin{pmatrix} -0.124 \\ 0.430 \\ -0.200 \\ 0.329 \end{pmatrix}$$

$$v_{\text{the}} = \begin{pmatrix} 0.234 \\ 0.266 \\ 0.239 \\ -0.199 \end{pmatrix} \quad v_{\text{language}} = \begin{pmatrix} 0.290 \\ -0.441 \\ 0.762 \\ 0.982 \end{pmatrix}$$

Similar words are “**nearby in the vector space**”



(Bandyopadhyay et al. 2022)
Antonie Wind
Go to Settings to a

Why word meaning in NLP models?

- With words, a feature is a word identity (= string)
 - Feature 5: `The previous word was “terrible”’
 - Requires **exact same word** to be in the training and testing set

“terrible” \neq “horrible”
- If we can represent word meaning in vectors:
 - The previous word was vector [35, 22, 17, ...]
 - Now in the test set we might see a similar vector [34, 21, 14, ...]
 - We can generalize to **similar but unseen** words!!!

Content

- Representing words
- Word2Vec
- Application of Word2Vec

How do we represent the meaning of a word?

Definition: **meaning** (Webster dictionary)

- the idea that is represented by a word, phrase, etc.
- the idea that a person wants to express by using words, signs, etc.
- the idea that is expressed in a work of writing, art, etc.

How do we have usable meaning in a computer?

Common solution: Use e.g. [WordNet](#), a thesaurus containing lists of **synonym sets** and **hypercnyms** (“is a” relationships).

e.g. synonym sets containing “good”:

```
from nltk.corpus import wordnet as wn
poses = { 'n':'noun', 'v':'verb', 's':'adj (s)', 'a':'adj', 'r':'adv'}
for synset in wn.synsets("good"):
    print("{}: {}".format(poses[synset.pos()],
        ", ".join([l.name() for l in synset.lemmas()])))
```

```
noun: good
noun: good, goodness
noun: good, goodness
noun: commodity, trade_good, good
adj: good
adj (sat): full, good
adj: good
adj (sat): estimable, good, honorable, respectable
adj (sat): beneficial, good
adj (sat): good
adj (sat): good, just, upright
...
adverb: well, good
adverb: thoroughly, soundly, good
```

e.g. hypernyms of “panda”:

```
from nltk.corpus import wordnet as wn
panda = wn.synset("panda.n.01")
hyper = lambda s: s.hypernyms()
list(pandaclosure(hyper))
```

```
[Synset('procyonid.n.01'),
Synset('carnivore.n.01'),
Synset('placental.n.01'),
Synset('mammal.n.01'),
Synset('vertebrate.n.01'),
Synset('chordate.n.01'),
Synset('animal.n.01'),
Synset('organism.n.01'),
Synset('living_thing.n.01'),
Synset('whole.n.02'),
Synset('object.n.01'),
Synset('physical_entity.n.01'),
Synset('entity.n.01')]
```

Problems with resources like WordNet

- Great as a resource but missing nuance
 - e.g. “proficient” is listed as a synonym for “good”. This is only correct in some contexts.
- Missing new meanings of words
 - e.g., wicked, badass, nifty, wizard, genius, ninja, bombest
 - Impossible to keep up-to-date!
- Subjective
- Requires human labor to create and adapt
- Can’t compute accurate word similarity →

Representing words as discrete symbols

In traditional NLP, we regard words as discrete symbols:
`hotel`, `conference`, `motel` – a **localist** representation

Means one 1, the rest 0s

Words can be represented by **one-hot** vectors:

`motel = [0 0 0 0 0 0 0 0 0 1 0 0 0]`

`hotel = [0 0 0 0 0 0 1 0 0 0 0 0 0]`

Vector dimension = number of words in vocabulary (e.g., 500,000)

Problem with words as discrete symbols

Example: in web search, if user searches for “Seattle motel”, we would like to match documents containing “Seattle hotel”.

But:

```
motel = [0 0 0 0 0 0 0 0 0 1 0 0 0 0]  
hotel = [0 0 0 0 0 0 1 0 0 0 0 0 0 0]
```

These two vectors are **orthogonal**.

There is no natural notion of **similarity** for one-hot vectors!

Solution:

- Could try to rely on WordNet’s list of synonyms to get similarity?
 - But it is well-known to fail badly: incompleteness, etc.
- **Instead: learn to encode similarity in the vectors themselves**

Representing words by their context

- Distributional semantics: A word's meaning is given by the words that frequently appear close-by



- "You shall know a word by the company it keeps" (J. R. Firth 1957: 11)
- "Words which frequently appear in similar contexts have similar meaning"
- One of the most successful ideas of modern statistical NLP!

- When a word w appears in a text, its **context** is the set of words that appear nearby (within a fixed-size window)
- Use the many contexts of w to build up a representation of w

...government debt problems turning into banking crises as happened in 2009...

...saying that Europe needs unified banking regulation to replace the hodgepodge...

...India has just given its banking system a shot in the arm...

These **context words** will represent **banking**

Distributed representation

- Vector representation that encodes information about the distribution of contexts a word appears in
- Words that appear in similar contexts have similar representations
- We have several different ways we can encode the notion of “context.

Term-document matrix


- Context = appearing **in the same document**.

	Hamlet	Macbeth	Romeo & Juliet	Richard III	Julius Caesar	Tempest	Othello	King Lear
knife	1	1	4	2		2		10
dog				6	12	2		
sword	2	2	7	5		5		17
love	64		135	63		12		48
like	75	38	34	36	34	41	27	44

Measuring similarity

$$\cos(x, y) = \frac{\sum_{i=1}^F x_i y_i}{\sqrt{\sum_{i=1}^F x_i^2} \sqrt{\sum_{i=1}^F y_i^2}}$$

- We can calculate the cosine similarity of two vectors to judge the degree of their similarity [Salton 1971]
- Cosine similarity measures their orientation
- A common similarity metric: **cosine** of the angle between the two vectors (the larger, the more similar the two vectors are)



$\cos(\text{knife}, \text{knife})$	1
$\cos(\text{knife}, \text{dog})$	0.11
$\cos(\text{knife}, \text{sword})$	0.99
$\cos(\text{knife}, \text{love})$	0.65
$\cos(\text{knife}, \text{like})$	0.61

Sparse vs dense vectors

- The vectors in the word-word occurrence matrix are
 - **Long**: vocabulary size
 - **Sparse**: most are 0's
- Alternative: we want to represent words as **short** (50-300 dimensional) & **dense** (real-valued) vectors
 - The basis for modern NLP systems

$$v_{\text{cat}} = \begin{pmatrix} -0.224 \\ 0.130 \\ -0.290 \\ 0.276 \end{pmatrix} \quad v_{\text{dog}} = \begin{pmatrix} -0.124 \\ 0.430 \\ -0.200 \\ 0.329 \end{pmatrix}$$

$$v_{\text{the}} = \begin{pmatrix} 0.234 \\ 0.266 \\ 0.239 \\ -0.199 \end{pmatrix} \quad v_{\text{language}} = \begin{pmatrix} 0.290 \\ -0.441 \\ 0.762 \\ 0.982 \end{pmatrix}$$

Word vectors

- We will build a **dense vector** for each word, chosen so that it is similar to vectors of words that appear in similar contexts


$$\text{banking} = \begin{pmatrix} 0.286 \\ 0.792 \\ -0.177 \\ -0.107 \\ 0.109 \\ -0.542 \\ 0.349 \\ 0.271 \end{pmatrix}$$

- Note: **word vectors** are sometimes called **word embeddings** or **word representations**. They are a **distributed** representation.


Word embeddings: idea

- **Idea:** We have to put information about contexts into word vectors.
- **How:** Learn word vectors by teaching them to predict contexts
- **Prior work:**
 - Learning representations by back-propagating errors. (Rumelhart et al., 1986)
 - A neural probabilistic language model (**Bengio et al., 2003**)
 - NLP (almost) from Scratch (Collobert & Weston, 2008)
 - A recent, even simpler and faster model: word2vec (Mikolov et al. 2013)

Language model based on neural networks (Bengio et al., 2003)




Embedding matrix



Word embeddings: similarity

- Hope to have similar words nearby



Word embeddings: relationships

- Hope to preserve some language structure (relationships between words).

Relationship	Example 1	Example 2	Example 3
France - Paris	Italy: Rome	Japan: Tokyo	Florida: Tallahassee
big - bigger	small: larger	cold: colder	quick: quicker
Miami - Florida	Baltimore: Maryland	Dallas: Texas	Kona: Hawaii
Einstein - scientist	Messi: midfielder	Mozart: violinist	Picasso: painter
Sarkozy - France	Berlusconi: Italy	Merkel: Germany	Koizumi: Japan
copper - Cu	zinc: Zn	gold: Au	uranium: plutonium
Berlusconi - Silvio	Sarkozy: Nicolas	Putin: Medvedev	Obama: Barack
Microsoft - Windows	Google: Android	IBM: Linux	Apple: iPhone
Microsoft - Ballmer	Google: Yahoo	IBM: McNealy	Apple: Jobs
Japan - sushi	Germany: bratwurst	France: tapas	USA: pizza

Word embeddings: questions

- How big should the embedding space be?
 - Trade-offs like any other machine learning problem – greater capacity versus efficiency and overfitting.
- How do we find W matrix (vectors)?
 - Often as part of a prediction or classification task involving neighboring words.

Word2vec: Overview


- Word2vec (Mikolov et al. 2013) is a framework for learning word vectors

Idea:

- We have a large corpus of text
- Every word in a fixed vocabulary is represented by a vector
- Go through each position t in the text, which has a center word c and context (“outside”) words o
- Use the similarity of the word vectors for c and o to calculate the probability of o given c (or vice versa)
- Keep adjusting the word vectors to maximize this probability


Word2Vec Overview

- Example windows and process for computing $P(W_{t+j} | w_t)$




Word2Vec Overview

- Example windows and process for computing $P(W_{t+j} | W_t)$



word2vec

- Predict words using context
- Two versions: CBOW (continuous bag of words) and Skip-gram




Skip gram/CBOW intuition

- Similar “contexts” (that is, what words are likely to appear around them), lead to similar embeddings for two words.
- One way for the network to output similar context predictions for these two words is if *the word vectors are similar*. So, if two words have similar contexts, then the network is motivated to learn similar word vectors for these two words!


CBOW


- Bag of words
 - Gets rid of word order. Used in discrete case using counts of words that appear.
- CBOW
 - Takes vector embeddings of n words before target and n words after and adds them (as vectors).
 - Also removes word order, but the vector sum is meaningful enough to deduce missing word.





Word2vec – Continuous Bag of Word


- E.g. “The cat sat on floor”
 - Window size = 2




















We can consider either W or W' as the word's representation. Or even take the average.

Details

$$\text{Loss} = J(\theta) = -\frac{1}{T} \log L(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} \log P(w_{t+j} | w_t, \theta) = \frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} J_{t,j}(\theta).$$

... I saw a cute grey cat playing in the garden ...


with the central word *cat*, and four context words. Since we are going to look at just one step, we will pick only one of the context words; for example, let's take *cute*. Then the loss term for the central word *cat* and the context word *cute* is:

$$J_{t,j}(\theta) = -\log P(\text{cute} | \text{cat}) = -\log \frac{\exp u_{\text{cute}}^T v_{\text{cat}}}{\sum_{w \in \text{Voc}} \exp u_w^T v_{\text{cat}}} = -u_{\text{cute}}^T v_{\text{cat}} + \log \sum_{w \in \text{Voc}} \exp u_w^T v_{\text{cat}}.$$

Note which parameters are present at this step:

- from vectors for central words, only v_{cat} ;
- from vectors for context words, all u_w (for all words in the vocabulary).

1. Take dot product of v_{cat} with all u



2. exp

The diagram shows the result of step 1 for several words w : $u_{w1}^T v_{cat}$, $u_{w3}^T v_{cat}$, ..., $u_{cute}^T v_{cat}$, ..., $u_{wn}^T v_{cat}$. Each is followed by an arrow pointing to its exponential value: $\exp(u_{w1}^T v_{cat})$, $\exp(u_{w3}^T v_{cat})$, ..., $\exp(u_{cute}^T v_{cat})$, ..., $\exp(u_{wn}^T v_{cat})$. These values are summed to produce the final result: $\sum_{w \in V} \exp(u_w^T v_{cat})$.

3. sum all


4. get loss (for this one step)

$$J_{t,j}(\theta) = -\underbrace{u_{cute}^T v_{cat}}_{1} + \log \underbrace{\sum_{w \in V} \exp(u_w^T v_{cat})}_{2}$$

5. evaluate the gradient,
make an update

$$v_{cat} := v_{cat} - \alpha \frac{\partial J_{t,j}(\theta)}{\partial v_{cat}}$$
$$u_w := u_w - \alpha \frac{\partial J_{t,j}(\theta)}{\partial u_w} \quad \forall w \in V$$

By making an update to minimize $J_{t,j}(\theta)$, we force the parameters to increase similarity (dot product) of v_{cat} and u_{cute} and, at the same time, to decrease similarity between v_{cat} and u_w for all other words w in the vocabulary.



Some interesting results

Word Analogies

Test for linear relationships, examined by Mikolov et al. (2014)

a:b :: c:?



$$d = \arg \max_x \frac{(w_b - w_a + w_c)^T w_x}{\|w_b - w_a + w_c\|}$$


man:woman :: king:?

+ king [0.30 0.70]


- man [0.20 0.20]

+ woman [0.60 0.30]

queen [0.70 0.80]



Word analogies




Word embeddings: properties

- Need to have a function $W(\text{word})$ that returns a vector encoding that word.
- Similarity of words corresponds to nearby vectors.
 - Director – chairman, scratched – scraped
- Relationships between words correspond to difference between vectors.
 - Big – bigger, small – smaller

Word embeddings: properties

- Relationships between words correspond to difference between vectors.




$$W(\text{"woman"}) - W(\text{"man"}) \simeq W(\text{"aunt"}) - W(\text{"uncle"})$$

$$W(\text{"woman"}) - W(\text{"man"}) \simeq W(\text{"queen"}) - W(\text{"king"})$$


Skip gram

- Skip gram – alternative to CBOW
 - Start with a single word embedding and try to predict the context words (surrounding words).
 - Skip-Gram works well with small datasets, and can better represent less frequent words.



Skip gram


- Map from center word to probability on surrounding words. One input/output unit below.
 - There is no activation function on the hidden layer neurons, but the output neurons use softmax.



Skip gram example

- Vocabulary of 10,000 words.
- Embedding vectors with 300 features.
- So the hidden layer is going to be represented by a weight matrix with 10,000 rows (multiply by vector on the left).

$$[0 \ 0 \ 0 \ 1 \ 0] \times \begin{bmatrix} 17 & 24 & 1 \\ 23 & 5 & 7 \\ 4 & 6 & 13 \\ 10 & 12 & 19 \\ 11 & 18 & 25 \end{bmatrix} = [10 \ 12 \ 19]$$



Word2vec shortcomings

- **Problem:** 10,000 words and 300 dim embedding gives a large parameter space to learn. And 10K words is minimal for real applications.
- **Slow to train**, and need lots of data, particularly to learn uncommon words.
- Both of these layers would have a weight matrix with $300 \times 10,000 = 3$ million weights each

Faster Training: Negative Sampling


Dot product of v_{cat} :

- with u_{cute} - increase,
- with all other u - decrease




Dot product of v_{cat} :

- with u_{cute} - increase,
- with a subset of other u - decrease



Negative samples: randomly selected K words



Parameters to be updated:


- v_{cat}
 - u_w for all w in the vocabulary
- $|V| + 1$ vectors

Parameters to be updated:

- v_{cat}
 - u_{cute} and u_w for w in K negative examples
- $K + 2$ vectors

Defining a new learning problem

I want a glass of orange juice to go along with my cereal.




$k = 5-20$ smaller datasets

$k = 2-5$ larger dataset

Activate Wind
+ Cereals

Word2vec improvements: selective updates



Word2vec improvements: selective updates

- Selecting 5-20 words works well for smaller datasets, and you can get away with only 2-5 words for large datasets.
- Updating the weights for our positive word (“*jump*”), plus the weights for 5 other words that we want to output 0. That’s a total of 6 output neurons, and 1,800 weight values total. That’s only **0.06% of the 3M weights** in the output layer!

The GloVe Model (2014)

- GloVe observes that ratios of word-word co-occurrence probabilities have the potential for encoding some form of meaning. To consider the co-occurrence probabilities for target words **ice** and **steam** with various probe words from the vocabulary:
 - As one might expect, **ice** co-occurs more frequently with **solid** than it does with **gas**, whereas **steam** co-occurs more frequently with **gas** than it does with **solid**.
 - Both words co-occur with their shared property **water** frequently, and both co-occur with the unrelated word **fashion** infrequently.

Word2Vec vs. GloVe


- The advantage of GloVe is that, unlike Word2vec, GloVe does not rely just on local statistics (local context information of words), but incorporates global statistics (word co-occurrence) to obtain word vectors.

FastText vs. Word2Vec

- Fasttext invented by Facebook (2016)
- Fasttext solving unknown word problem.
- FastText operates at a more granular level with character n-grams. Where words are represented by the sum of the character n-gram vectors.

Word embedding applications


- The use of word representations... has become a key “secret sauce” for the success of many NLP systems in recent years, across tasks including named entity recognition, part-of-speech tagging, parsing, and semantic role labeling. ([Luong et al. \(2013\)](#))
- Learning a good representation on a task A and then using it on a task B is one of the major tricks in the Deep Learning toolbox.
 - Pretraining, transfer learning, and multi-task learning, **Large Language Models**.
 - Can allow the representation to learn from more than one kind of data.
- StarSpace: Embed All The Things (Ledell Wu et al, 2017)




W and *F* learn to perform task A. Later, *G* can learn to perform B based on *W*.

Word embedding applications

- Can apply to get a joint embedding of words and images or other multi-modal data sets.
- New classes map near similar existing classes: e.g., if ‘cat’ is unknown, cat images map near dog.



(Socher *et al.* (2013b))



Conclusion

- Representing words
- Word2vec
 - Skip-gram
 - CBOW
- Application of word2vec

References

- Slide of NLP Course, Stanford, 2022
- Slide of NLP Course, Princeton, 2023
- Other documents on Internet.